

✓ Preparation

```

!git clone https://github.com/mlvlab/ProMetaR.git
%cd ProMetaR/

!git clone https://github.com/KaiyangZhou/Dassl.pytorch.git
%cd Dassl.pytorch/

# Install dependencies
!pip install -r requirements.txt
!cp -r dassl ../
# Install this library (no need to re-build if the source code is modified)
# !python setup.py develop
%cd ..

!pip install -r requirements.txt

%mkdir outputs
%mkdir data

%cd data
%mkdir eurosat
!wget http://madm.dfki.de/files/sentinel/EuroSAT.zip EuroSAT.zip

!unzip -o EuroSAT.zip -d eurosat/
%cd eurosat
!gdown 1Ip7yaCWF0ea0FUGga0lUdVi_DDQth1o

%cd ../..

import os.path as osp
from collections import OrderedDict
import math
import torch
import torch.nn as nn
from torch.nn import functional as F
from torch.cuda.amp import GradScaler, autocast
from PIL import Image
import torchvision.transforms as transforms
import torch
from clip import clip
from clip.simple_tokenizer import SimpleTokenizer as _Tokenizer
import time
from tqdm import tqdm
import datetime
import argparse
from dassl.utils import setup_logger, set_random_seed, collect_env_info
from dassl.config import get_cfg_default
from dassl.engine import build_trainer
from dassl.engine import TRAINER_REGISTRY, TrainerX
from dassl.metrics import compute_accuracy
from dassl.utils import load_pretrained_weights, load_checkpoint
from dassl.optim import build_optimizer, build_lr_scheduler

# custom
import datasets.oxford_pets
import datasets.oxford_flowers
import datasets.fgvc_aircraft
import datasets.dtd
import datasets.eurosat
import datasets.stanford_cars
import datasets.food101
import datasets.sun397
import datasets.caltech101
import datasets.ucf101
import datasets.imagenet
import datasets.imagenet_sketch
import datasets.imagenetv2
import datasets.imagenet_a
import datasets.imagenet_r

def print_args(args, cfg):
    ... print("*****")
    ... print("** Arguments **")
    ... print("*****")

```

```

... optkeys = list(args.__dict__.keys())
... optkeys.sort()
... for key in optkeys:
...     print("{}: {}".format(key, args.__dict__[key]))
... print("*****")
... print("** Config **")
... print("*****")
... print(cfg)

def reset_cfg(cfg, args):
... if args.root:
...     cfg.DATASET.ROOT = args.root
... if args.output_dir:
...     cfg.OUTPUT_DIR = args.output_dir
... if args.seed:
...     cfg.SEED = args.seed
... if args.trainer:
...     cfg.TRAINER.NAME = args.trainer
... cfg.DATASET.NUM_SHOTS = 16
... cfg.DATASET.SUBSAMPLE_CLASSES = args.subsample_classes
... cfg.DATALOADER.TRAIN_X.BATCH_SIZE = args.train_batch_size
... cfg.OPTIM.MAX_EPOCH = args.epoch

def extend_cfg(cfg):
... """
... Add new config variables.
... """
... from yacs.config import CfgNode as CN
... cfg.TRAINER.COOP = CN()
... cfg.TRAINER.COOP.N_CTX = 16 # number of context vectors
... cfg.TRAINER.COOP.CSC = False # class-specific context
... cfg.TRAINER.COOP.CTX_INIT = "" # initialization words
... cfg.TRAINER.COOP.PREC = "fp16" # fp16, fp32, amp
... cfg.TRAINER.COOP.CLASS_TOKEN_POSITION = "end" # 'middle' or 'end' or 'front'
... cfg.TRAINER.COCOOP = CN()
... cfg.TRAINER.COCOOP.N_CTX = 4 # number of context vectors
... cfg.TRAINER.COCOOP.CTX_INIT = "a photo of a" # initialization words
... cfg.TRAINER.COCOOP.PREC = "fp16" # fp16, fp32, amp
... cfg.TRAINER.PROMETAR = CN()
... cfg.TRAINER.PROMETAR.N_CTX_VISION = 4 # number of context vectors at the vision branch
... cfg.TRAINER.PROMETAR.N_CTX_TEXT = 4 # number of context vectors at the language branch
... cfg.TRAINER.PROMETAR.CTX_INIT = "a photo of a" # initialization words
... cfg.TRAINER.PROMETAR.PREC = "fp16" # fp16, fp32, amp
... cfg.TRAINER.PROMETAR.PROMPT_DEPTH_VISION = 9 # Max 12, minimum 0, for 0 it will be using shallow IVLP prompting (
... cfg.TRAINER.PROMETAR.PROMPT_DEPTH_TEXT = 9 # Max 12, minimum 0, for 0 it will be using shallow IVLP prompting (J=
... cfg.DATASET.SUBSAMPLE_CLASSES = "all" # all, base or new
... cfg.TRAINER.PROMETAR.ADAPT_LR = 0.0005
... cfg.TRAINER.PROMETAR.LR_RATIO = 0.0005
... cfg.TRAINER.PROMETAR.FAST_ADAPTATION = False
... cfg.TRAINER.PROMETAR.MIXUP_ALPHA = 0.5
... cfg.TRAINER.PROMETAR.MIXUP_BETA = 0.5
... cfg.TRAINER.PROMETAR.DIM_RATE=8
... cfg.OPTIM_VNET = CN()
... cfg.OPTIM_VNET.NAME = "adam"
... cfg.OPTIM_VNET.LR = 0.0003
... cfg.OPTIM_VNET.WEIGHT_DECAY = 5e-4
... cfg.OPTIM_VNET.MOMENTUM = 0.9
... cfg.OPTIM_VNET.SGD_DAMPNING = 0
... cfg.OPTIM_VNET.SGD_NESTEROV = False
... cfg.OPTIM_VNET.RMSPROP_ALPHA = 0.99
... cfg.OPTIM_VNET.ADAM_BETA1 = 0.9
... cfg.OPTIM_VNET.ADAM_BETA2 = 0.999
... cfg.OPTIM_VNET.STAGED_LR = False
... cfg.OPTIM_VNET.NEW_LAYERS = ()
... cfg.OPTIM_VNET.BASE_LR_MULT = 0.1
... # Learning rate scheduler
... cfg.OPTIM_VNET.LR_SCHEDULER = "single_step"
... # -1 or 0 means the stepsize is equal to max_epoch
... cfg.OPTIM_VNET.STEPSIZE = (-1, )
... cfg.OPTIM_VNET.GAMMA = 0.1
... cfg.OPTIM_VNET.MAX_EPOCH = 10
... # Set WARMUP_EPOCH larger than 0 to activate warmup training
... cfg.OPTIM_VNET.WARMUP_EPOCH = -1
... # Either linear or constant
... cfg.OPTIM_VNET.WARMUP_TYPE = "linear"
... # Constant learning rate when type=constant
... cfg.OPTIM_VNET.WARMUP_CONS_LR = 1e-5
... # Minimum learning rate when type=linear

```

```

... cfg.OPTIM_VNET.WARMUP_MIN_LR = 1e-5
... # Recount epoch for the next scheduler (last_epoch=-1)
... # Otherwise last_epoch=warmup_epoch
... cfg.OPTIM_VNET.WARMUP_RECOUNT = True

def setup_cfg(args):
... cfg = get_cfg_default()
... extend_cfg(cfg)
... # 1. From the dataset config file
... if args.dataset_config_file:
...     cfg.merge_from_file(args.dataset_config_file)
... # 2. From the method config file
... if args.config_file:
...     cfg.merge_from_file(args.config_file)
... # 3. From input arguments
... reset_cfg(cfg, args)
... cfg.freeze()
... return cfg

_tokenizer = _Tokenizer()

def load_clip_to_cpu(cfg): # Load CLIP
... backbone_name = cfg.MODEL.BACKBONE.NAME
... url = clip._MODELS[backbone_name]
... model_path = clip._download(url)

... try:
...     # loading JIT archive
...     model = torch.jit.load(model_path, map_location="cpu").eval()
...     state_dict = None

... except RuntimeError:
...     state_dict = torch.load(model_path, map_location="cpu")

... if cfg.TRAINER.NAME == "":
...     design_trainer = "CoOp"
... else:
...     design_trainer = cfg.TRAINER.NAME
... design_details = {"trainer": design_trainer,
...                   "vision_depth": 0,
...                   "language_depth": 0, "vision_ctx": 0,
...                   "language_ctx": 0}
... model = clip.build_model(state_dict or model.state_dict(), design_details)

... return model

from dassl.config import get_cfg_default
cfg = get_cfg_default()
cfg.MODEL.BACKBONE.NAME = "ViT-B/16" # Set the vision encoder backbone of CLIP to ViT.
clip_model = load_clip_to_cpu(cfg)

class TextEncoder(nn.Module):
... def __init__(self, clip_model): # 초기화 하는 함수
...     super().__init__()
...     self.transformer = clip_model.transformer
...     self.positional_embedding = clip_model.positional_embedding
...     self.ln_final = clip_model.ln_final
...     self.text_projection = clip_model.text_projection
...     self.dtype = clip_model.dtype

... def forward(self, prompts, tokenized_prompts): # 모델 호출
...     x = prompts + self.positional_embedding.type(self.dtype)
...     x = x.permute(1, 0, 2) # NLD -> LND
...     x = self.transformer(x)
...     x = x.permute(1, 0, 2) # LND -> NLD
...     x = self.ln_final(x).type(self.dtype)

...     # x.shape = [batch_size, n_ctx, transformer.width]
...     # take features from the eot embedding (eot_token is the highest number in each sequence)
...     x = x[torch.arange(x.shape[0]), tokenized_prompts.argmax(dim=-1)] @ self.text_projection

...     return x

@TRAINER_REGISTRY.register(force=True)
class CoCoOp(TrainerX):

```

```

def check_cfg(self, cfg):
    assert cfg.TRAINER.COCOOP.PREC in ["fp16", "fp32", "amp"]

def build_model(self):
    cfg = self.cfg
    classnames = self.dm.dataset.classnames
    print(f"Loading CLIP (backbone: {cfg.MODEL.BACKBONE.NAME})")
    clip_model = load_clip_to_cpu(cfg)

    if cfg.TRAINER.COCOOP.PREC == "fp32" or cfg.TRAINER.COCOOP.PREC == "amp":
        # CLIP's default precision is fp16
        clip_model.float()

    print("Building custom CLIP")
    self.model = CoCoOpCustomCLIP(cfg, classnames, clip_model)

    print("Turning off gradients in both the image and the text encoder")
    name_to_update = "prompt_learner"

    for name, param in self.model.named_parameters():
        if name_to_update not in name:
            param.requires_grad_(False)

    # Double check
    enabled = set()
    for name, param in self.model.named_parameters():
        if param.requires_grad:
            enabled.add(name)
    print(f"Parameters to be updated: {enabled}")

    if cfg.MODEL.INIT_WEIGHTS:
        load_pretrained_weights(self.model.prompt_learner, cfg.MODEL.INIT_WEIGHTS)

    self.model.to(self.device)
    # NOTE: only give prompt_learner to the optimizer
    self.optim = build_optimizer(self.model.prompt_learner, cfg.OPTIM)
    self.sched = build_lr_scheduler(self.optim, cfg.OPTIM)
    self.register_model("prompt_learner", self.model.prompt_learner, self.optim, self.sched)

    self.scaler = GradScaler() if cfg.TRAINER.COCOOP.PREC == "amp" else None

    # Note that multi-gpu training could be slow because CLIP's size is
    # big, which slows down the copy operation in DataParallel
    device_count = torch.cuda.device_count()
    if device_count > 1:
        print(f"Multiple GPUs detected (n_gpus={device_count}), use all of them!")
        self.model = nn.DataParallel(self.model)

def before_train(self):
    directory = self.cfg.OUTPUT_DIR
    if self.cfg.RESUME:
        directory = self.cfg.RESUME
    self.start_epoch = self.resume_model_if_exist(directory)

    # Remember the starting time (for computing the elapsed time)
    self.time_start = time.time()

def forward_backward(self, batch):
    image, label = self.parse_batch_train(batch)

    model = self.model
    optim = self.optim
    scaler = self.scaler

    prec = self.cfg.TRAINER.COCOOP.PREC
    loss = model(image, label) # Input image 모델 통과
    optim.zero_grad()
    loss.backward() # Backward (역전파)
    optim.step() # 모델 parameter update

    loss_summary = {"loss": loss.item()}

    if (self.batch_idx + 1) == self.num_batches:
        self.update_lr()

    return loss_summary

def parse_batch_train(self, batch):

```

```

def parse_batch_train(self, batch):
    input = batch["img"]
    label = batch["label"]
    input = input.to(self.device)
    label = label.to(self.device)
    return input, label

def load_model(self, directory, epoch=None):
    if not directory:
        print("Note that load_model() is skipped as no pretrained model is given")
        return

    names = self.get_model_names()

    # By default, the best model is loaded
    model_file = "model-best.pth.tar"

    if epoch is not None:
        model_file = "model.pth.tar-" + str(epoch)

    for name in names:
        model_path = osp.join(directory, name, model_file)

        if not osp.exists(model_path):
            raise FileNotFoundError('Model not found at "{}".format(model_path))

        checkpoint = load_checkpoint(model_path)
        state_dict = checkpoint["state_dict"]
        epoch = checkpoint["epoch"]

        # Ignore fixed token vectors
        if "token_prefix" in state_dict:
            del state_dict["token_prefix"]

        if "token_suffix" in state_dict:
            del state_dict["token_suffix"]

        print("Loading weights to {} from {}".format(model_path, epoch))
        # set strict=False
        self._models[name].load_state_dict(state_dict, strict=False)

    def after_train(self):
        print("Finish training")

        do_test = not self.cfg.TEST.NO_TEST
        if do_test:
            if self.cfg.TEST.FINAL_MODEL == "best_val":
                print("Deploy the model with the best val performance")
                self.load_model(self.output_dir)
            else:
                print("Deploy the last-epoch model")
            acc = self.test()

        # Show elapsed time
        elapsed = round(time.time() - self.time_start)
        elapsed = str(datetime.timedelta(seconds=elapsed))
        print(f"Elapsed: {elapsed}")

        # Close writer
        self.close_writer()
        return acc

    def train(self):
        """Generic training loops."""
        self.before_train()
        for self.epoch in range(self.start_epoch, self.max_epoch):
            self.before_epoch()
            self.run_epoch()
            self.after_epoch()
        acc = self.after_train()
        return acc

parser = argparse.ArgumentParser()
parser.add_argument("--root", type=str, default="data/", help="path to dataset")
parser.add_argument("--output-dir", type=str, default="outputs/cocoop3", help="output directory")
parser.add_argument(
    "--seed", type=int, default=1, help="only positive value enables a fixed seed"
)
parser.add_argument(

```

```

parser.add_argument(
    "--config-file", type=str, default="configs/trainers/ProMetaR/vit_b16_c2_ep10_batch4_4ctx.yaml", help="path to c
)
parser.add_argument(
    "--dataset-config-file",
    type=str,
    default="configs/datasets/eurosat.yaml",
    help="path to config file for dataset setup",
)
parser.add_argument("--trainer", type=str, default="CoOp", help="name of trainer")
parser.add_argument("--eval-only", action="store_true", help="evaluation only")
parser.add_argument(
    "--model-dir",
    type=str,
    default="",
    help="load model from this directory for eval-only mode",
)
parser.add_argument("--train-batch-size", type=int, default=4)
parser.add_argument("--epoch", type=int, default=10)
parser.add_argument("--subsample-classes", type=str, default="base")
parser.add_argument(
    "--load-epoch", type=int, default=0, help="load model weights at this epoch for evaluation"
)
args = parser.parse_args([])

def main(args):
    cfg = setup_cfg(args)
    if cfg.SEED >= 0:
        set_random_seed(cfg.SEED)

    if torch.cuda.is_available() and cfg.USE_CUDA:
        torch.backends.cudnn.benchmark = True

    trainer = build_trainer(cfg)
    if args.eval_only:
        trainer.load_model(args.model_dir, epoch=args.load_epoch)
        acc = trainer.test()
        return acc

    acc = trainer.train()
    return acc

```



```

inflating: eurosat/2750/PermanentCrop/PermanentCrop_579.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_1025.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_2409.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_853.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_421.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_386.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_2068.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_882.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_357.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_1.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_65.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_736.jpg
/content/ProMetaR/data/eurosat
Downloading...
From: https://drive.google.com/uc?id=1Ip7yaCWFj0ea0FUGGa0lUdVi\_DD0th1o
To: /content/ProMetaR/data/eurosat/split_zhou_EuroSAT.json
100% 3.01M/3.01M [00:00<00:00, 23.6MB/s]
/content/ProMetaR

```

```
!ls data/eurosat/2750
```

```

AnnualCrop  HerbaceousVegetation  Industrial  PermanentCrop  River
Forest      Highway              Pasture     Residential    SeaLake

```

✓ Use another architecture for CoCoOp's Meta Net - deeper

Key changes made to the meta network:

Original architecture:

- Single linear transformation with ReLU:
 - Linear(vis_dim → vis_dim // 16) + ReLU
 - Linear(vis_dim // 16 → ctx_dim)

New deeper architecture:

- Four-layer deep network:
 - First layer: Linear(vis_dim → hidden_dim) + ReLU
 - Second layer: Linear(hidden_dim → hidden_dim) + ReLU
 - Third layer: Linear(hidden_dim → hidden_dim) + ReLU
 - Fourth layer: Linear(hidden_dim → ctx_dim)

Benefits of deeper architecture:

- Enhanced feature extraction capability
- Increased model capacity
- More complex function approximation
- Hierarchical representation learning

Changes in implementation:

- Maintained the same hidden dimension (vis_dim // 16)
- Added two additional hidden layers with ReLU activation
- Kept the final output dimension unchanged (ctx_dim)
- No changes to the bias calculation process

This modification focuses purely on increasing the depth of the network while maintaining the same input and output dimensions, allowing the model to learn more complex transformations between the visual and context features.

```
import torch.nn as nn
```

```

class CoCoOpPromptLearner(nn.Module):
    def __init__(self, cfg, classnames, clip_model):
        super().__init__()
        n_cls = len(classnames)
        n_ctx = cfg.TRAINER.COOCOOP.N_CTX
        ctx_init = cfg.TRAINER.COOCOOP.CTX_INIT
        dtype = clip_model.dtype
        ctx_dim = clip_model.ln_final.weight.shape[0]
        vis_dim = clip_model.visual.output_dim
        clip_imsize = clip_model.visual.input_resolution
        cfg_imsize = cfg.INPUT.SIZE[0]
        assert cfg_imsize == clip_imsize, f"cfg_imsize ({cfg_imsize}) must equal to clip_imsize ({clip_imsize})"

```

```

if ctx_init:
    # use given words to initialize context vectors
    ctx_init = ctx_init.replace("_", " ")
    n_ctx = len(ctx_init.split(" "))
    prompt = clip.tokenize(ctx_init)
    with torch.no_grad():
        embedding = clip_model.token_embedding(prompt).type(dtype)
    ctx_vectors = embedding[0, 1: 1 + n_ctx, :]
    prompt_prefix = ctx_init
else:
    # random initialization
    ctx_vectors = torch.empty(n_ctx, ctx_dim, dtype=dtype)
    nn.init.normal_(ctx_vectors, std=0.02)
    prompt_prefix = " ".join(["X"] * n_ctx)

print(f'Initial context: "{prompt_prefix}"')
print(f"Number of context words (tokens): {n_ctx}")

self.ctx = nn.Parameter(ctx_vectors) # Wrap the initialized prompts above as parameters to make them trainable

### Tokenize ###
classnames = [name.replace("_", " ") for name in classnames] # 예) "Forest"
name_lens = [len(_tokenizer.encode(name)) for name in classnames]
prompts = [prompt_prefix + " " + name + "." for name in classnames] # 예) "A photo of Forest."

tokenized_prompts = torch.cat([clip.tokenize(p) for p in prompts]) # 예) [49406, 320, 1125, 539...]

# #####
# ##### Define Meta Net #####
# self.meta_net = nn.Sequential(OrderedDict([
#     # ("linear1", "fill in here"(vis_dim, vis_dim // 16)),
#     ("linear1", nn.Linear(vis_dim, vis_dim // 16)),
#     ("relu", nn.ReLU(inplace=True)),
#     ("linear2", nn.Linear(vis_dim // 16, ctx_dim))
# ]))
# #####

#####
##### Define Meta Net #####
hidden_dim = vis_dim // 16
self.meta_net = nn.Sequential(OrderedDict([
    ("layer1", nn.Sequential(
        nn.Linear(vis_dim, hidden_dim),
        nn.ReLU(inplace=True),
    )),
    ("layer2", nn.Sequential(
        nn.Linear(hidden_dim, hidden_dim),
        nn.ReLU(inplace=True),
    )),
    ("layer3", nn.Sequential(
        nn.Linear(hidden_dim, hidden_dim),
        nn.ReLU(inplace=True),
    )),
    ("layer4", nn.Linear(hidden_dim, ctx_dim))
]))
#####

if cfg.TRAINER.COCOP.PREC == "fp16":
    self.meta_net.half()

with torch.no_grad():
    embedding = clip_model.token_embedding(tokenized_prompts).type(dtype)

# These token vectors will be saved when in save_model(),
# but they should be ignored in load_model() as we want to use
# those computed using the current class names
self.register_buffer("token_prefix", embedding[:, :1, :]) # SOS
self.register_buffer("token_suffix", embedding[:, 1 + n_ctx:, :]) # CLS, EOS
self.n_cls = n_cls
self.n_ctx = n_ctx
self.tokenized_prompts = tokenized_prompts # torch.Tensor
self.name_lens = name_lens

def construct_prompts(self, ctx, prefix, suffix, label=None):
    # dim0 is either batch_size (during training) or n_cls (during testing)

```



```

# ctx: context tokens, with shape of (dim0, n_ctx, ctx_dim)
# prefix: the sos token, with shape of (n_cls, 1, ctx_dim)
# suffix: remaining tokens, with shape of (n_cls, *, ctx_dim)

if label is not None:
    prefix = prefix[label]
    suffix = suffix[label]

prompts = torch.cat(
    [
        prefix, # (dim0, 1, dim)
        ctx, # (dim0, n_ctx, dim)
        suffix, # (dim0, *, dim)
    ],
    dim=1,
)

return prompts

def forward(self, im_features):
    prefix = self.token_prefix
    suffix = self.token_suffix
    ctx = self.ctx # (n_ctx, ctx_dim)

#####
#Image feature is given as input to meta network # (batch, ctx_dim)
bias = self.meta_net(im_features) # (batch, ctx_dim)
bias = bias.unsqueeze(1) # (batch, 1, ctx_dim)
ctx = ctx.unsqueeze(0) # (1, n_ctx, ctx_dim)
#ctx_shifted = ctx + " Add meta token to context token" # (batch, n_ctx, ctx_dim)
ctx_shifted = ctx + bias # (batch, n_ctx, ctx_dim)
#####
#####

# Use instance-conditioned context tokens for all classes
prompts = []
for ctx_shifted_i in ctx_shifted:
    ctx_i = ctx_shifted_i.unsqueeze(0).expand(self.n_cls, -1, -1)
    pts_i = self.construct_prompts(ctx_i, prefix, suffix) # (n_cls, n_tkn, ctx_dim)
    prompts.append(pts_i)
prompts = torch.stack(prompts)

return prompts

class CoCoOpCustomCLIP(nn.Module):
    def __init__(self, cfg, classnames, clip_model):
        super().__init__()
        self.prompt_learner = CoCoOpPromptLearner(cfg, classnames, clip_model)
        self.tokenized_prompts = self.prompt_learner.tokenized_prompts
        self.image_encoder = clip_model.visual
        self.text_encoder = TextEncoder(clip_model)
        self.logit_scale = clip_model.logit_scale
        self.dtype = clip_model.dtype

    def forward(self, image, label=None):
        tokenized_prompts = self.tokenized_prompts
        logit_scale = self.logit_scale.exp()

        image_features = self.image_encoder(image.type(self.dtype))
        image_features = image_features / image_features.norm(dim=-1, keepdim=True)

#####
prompts = self.prompt_learner(image_features)
#####

logits = []
for pts_i, imf_i in zip(prompts, image_features):
    text_features = self.text_encoder(pts_i, tokenized_prompts)
    text_features = text_features / text_features.norm(dim=-1, keepdim=True)
    l_i = logit_scale * imf_i @ text_features.t()
    logits.append(l_i)

```

```

logits = torch.stack(logits)

if self.prompt_learner.training:
    return F.cross_entropy(logits, label)

return logits

```

▼ Training

- Train modified version of CoCoOp (deeper) on the EuroSAT dataset.

Train on the Base Classes Train split and evaluate accuracy on the Base Classes Test split.

```

args.trainer = "CoCoOp"
args.train_batch_size = 4
args.epoch = 100
args.output_dir = "outputs/cocoop"

```

```

args.subsample_classes = "base"
args.eval_only = False
cocoop_base_acc = main(args)

```

```

epoch [55/100] batch [20/20] time 0.096 (0.125) data 0.000 (0.016) loss 0.1265 (0.2092) lr 1.0933e-03 eta 0:01:52
epoch [56/100] batch [20/20] time 0.130 (0.141) data 0.000 (0.023) loss 0.2367 (0.3348) lr 1.0545e-03 eta 0:02:03
epoch [57/100] batch [20/20] time 0.152 (0.197) data 0.000 (0.032) loss 0.1893 (0.2059) lr 1.0158e-03 eta 0:02:49
epoch [58/100] batch [20/20] time 0.093 (0.127) data 0.000 (0.020) loss 0.2300 (0.3090) lr 9.7732e-04 eta 0:01:46
epoch [59/100] batch [20/20] time 0.097 (0.124) data 0.000 (0.016) loss 0.0651 (0.2129) lr 9.3914e-04 eta 0:01:41
epoch [60/100] batch [20/20] time 0.094 (0.125) data 0.000 (0.017) loss 0.0238 (0.2052) lr 9.0126e-04 eta 0:01:40
epoch [61/100] batch [20/20] time 0.130 (0.139) data 0.000 (0.022) loss 0.4592 (0.2695) lr 8.6373e-04 eta 0:01:48
epoch [62/100] batch [20/20] time 0.152 (0.206) data 0.000 (0.032) loss 0.1094 (0.2459) lr 8.2658e-04 eta 0:02:36
epoch [63/100] batch [20/20] time 0.103 (0.126) data 0.000 (0.021) loss 0.0539 (0.2055) lr 7.8984e-04 eta 0:01:33
epoch [64/100] batch [20/20] time 0.109 (0.127) data 0.000 (0.022) loss 0.3464 (0.3186) lr 7.5357e-04 eta 0:01:31
epoch [65/100] batch [20/20] time 0.093 (0.124) data 0.000 (0.016) loss 0.1292 (0.2239) lr 7.1778e-04 eta 0:01:27
epoch [66/100] batch [20/20] time 0.125 (0.137) data 0.000 (0.018) loss 0.3057 (0.2014) lr 6.8251e-04 eta 0:01:33
epoch [67/100] batch [20/20] time 0.151 (0.213) data 0.000 (0.029) loss 0.0312 (0.2030) lr 6.4781e-04 eta 0:02:20
epoch [68/100] batch [20/20] time 0.093 (0.129) data 0.000 (0.018) loss 0.0652 (0.2042) lr 6.1370e-04 eta 0:01:22
epoch [69/100] batch [20/20] time 0.096 (0.127) data 0.000 (0.017) loss 0.0966 (0.2174) lr 5.8022e-04 eta 0:01:18
epoch [70/100] batch [20/20] time 0.098 (0.128) data 0.000 (0.020) loss 0.2649 (0.2203) lr 5.4740e-04 eta 0:01:17
epoch [71/100] batch [20/20] time 0.120 (0.143) data 0.000 (0.016) loss 0.1203 (0.2104) lr 5.1527e-04 eta 0:01:22
epoch [72/100] batch [20/20] time 0.136 (0.196) data 0.000 (0.030) loss 0.9736 (0.1950) lr 4.8387e-04 eta 0:01:49
epoch [73/100] batch [20/20] time 0.095 (0.127) data 0.000 (0.018) loss 0.2593 (0.1552) lr 4.5322e-04 eta 0:01:08
epoch [74/100] batch [20/20] time 0.095 (0.129) data 0.000 (0.017) loss 0.4309 (0.1589) lr 4.2336e-04 eta 0:01:07
epoch [75/100] batch [20/20] time 0.097 (0.129) data 0.000 (0.018) loss 0.1902 (0.2479) lr 3.9432e-04 eta 0:01:04
epoch [76/100] batch [20/20] time 0.121 (0.141) data 0.000 (0.016) loss 0.1500 (0.2145) lr 3.6612e-04 eta 0:01:07
epoch [77/100] batch [20/20] time 0.143 (0.195) data 0.000 (0.027) loss 0.0165 (0.1663) lr 3.3879e-04 eta 0:01:29
epoch [78/100] batch [20/20] time 0.097 (0.130) data 0.000 (0.024) loss 0.0516 (0.2440) lr 3.1236e-04 eta 0:00:57
epoch [79/100] batch [20/20] time 0.097 (0.131) data 0.000 (0.016) loss 0.0469 (0.1345) lr 2.8686e-04 eta 0:00:54
epoch [80/100] batch [20/20] time 0.099 (0.130) data 0.000 (0.026) loss 0.0911 (0.1605) lr 2.6231e-04 eta 0:00:52
epoch [81/100] batch [20/20] time 0.131 (0.143) data 0.000 (0.018) loss 0.0233 (0.1552) lr 2.3873e-04 eta 0:00:54
epoch [82/100] batch [20/20] time 0.160 (0.200) data 0.000 (0.029) loss 0.4319 (0.1434) lr 2.1615e-04 eta 0:01:12
epoch [83/100] batch [20/20] time 0.095 (0.128) data 0.000 (0.022) loss 0.5342 (0.1633) lr 1.9459e-04 eta 0:00:43
epoch [84/100] batch [20/20] time 0.097 (0.128) data 0.000 (0.021) loss 0.6577 (0.3381) lr 1.7407e-04 eta 0:00:40
epoch [85/100] batch [20/20] time 0.094 (0.128) data 0.000 (0.024) loss 0.3354 (0.2112) lr 1.5462e-04 eta 0:00:38
epoch [86/100] batch [20/20] time 0.130 (0.143) data 0.000 (0.017) loss 0.2001 (0.1630) lr 1.3624e-04 eta 0:00:39
epoch [87/100] batch [20/20] time 0.138 (0.195) data 0.000 (0.035) loss 0.1278 (0.1340) lr 1.1897e-04 eta 0:00:50
epoch [88/100] batch [20/20] time 0.094 (0.126) data 0.000 (0.018) loss 0.1560 (0.1252) lr 1.0281e-04 eta 0:00:30
epoch [89/100] batch [20/20] time 0.095 (0.125) data 0.000 (0.017) loss 0.5410 (0.2455) lr 8.7779e-05 eta 0:00:27
epoch [90/100] batch [20/20] time 0.094 (0.128) data 0.000 (0.021) loss 0.2766 (0.1934) lr 7.3899e-05 eta 0:00:25
epoch [91/100] batch [20/20] time 0.130 (0.142) data 0.000 (0.024) loss 0.0363 (0.1597) lr 6.1179e-05 eta 0:00:25
epoch [92/100] batch [20/20] time 0.163 (0.195) data 0.000 (0.033) loss 0.0415 (0.2217) lr 4.9633e-05 eta 0:00:31
epoch [93/100] batch [20/20] time 0.102 (0.129) data 0.000 (0.025) loss 0.2732 (0.2631) lr 3.9271e-05 eta 0:00:18
epoch [94/100] batch [20/20] time 0.092 (0.129) data 0.000 (0.016) loss 0.5952 (0.1980) lr 3.0104e-05 eta 0:00:15
epoch [95/100] batch [20/20] time 0.097 (0.126) data 0.000 (0.018) loss 0.1069 (0.1488) lr 2.2141e-05 eta 0:00:12
epoch [96/100] batch [20/20] time 0.132 (0.145) data 0.000 (0.019) loss 0.4573 (0.2035) lr 1.5390e-05 eta 0:00:11
epoch [97/100] batch [20/20] time 0.139 (0.198) data 0.000 (0.035) loss 0.0201 (0.2065) lr 9.8566e-06 eta 0:00:11
epoch [98/100] batch [20/20] time 0.105 (0.129) data 0.000 (0.019) loss 0.0546 (0.1570) lr 5.5475e-06 eta 0:00:05
epoch [99/100] batch [20/20] time 0.097 (0.126) data 0.000 (0.020) loss 0.0547 (0.1876) lr 2.4666e-06 eta 0:00:02
epoch [100/100] batch [20/20] time 0.093 (0.128) data 0.000 (0.018) loss 0.6436 (0.2344) lr 6.1680e-07 eta 0:00:00
Checkpoint saved to outputs/cocoop/prompt_learner/model.pth.tar-100
Finish training
Deploy the last-epoch model
Evaluate on the *test* set
100%|██████████| 42/42 [01:05<00:00, 1.55s/it]=> result
* total: 4,200
* correct: 3,655
* accuracy: 87.0%
* error: 13.0%
* macro_f1: 87.1%
Elapsed: 0:06:22

```

Accuracy on the New Classes

```
# Accuracy on the new classes.
args.model_dir = "outputs/cocoop"
args.output_dir = "outputs/cocoop/new_classes"
args.subsample_classes = "new"
args.load_epoch = 100
args.eval_only = True
coop_novel_acc = main(args)
```

↗ Loading trainer: CoCoOp
Loading dataset: EuroSAT
Reading split from /content/ProMetaR/data/eurosat/split_zhou_EuroSAT.json
Loading preprocessed few-shot data from /content/ProMetaR/data/eurosat/split_fewshot/shot_16-seed_1.pkl
SUBSAMPLE NEW CLASSES!
Building transform_train
+ random resized crop (size=(224, 224), scale=(0.08, 1.0))
+ random flip
+ to torch tensor of range [0, 1]
+ normalization (mean=[0.48145466, 0.4578275, 0.40821073], std=[0.26862954, 0.26130258, 0.27577711])
Building transform_test
+ resize the smaller edge to 224
+ 224x224 center crop
+ to torch tensor of range [0, 1]
+ normalization (mean=[0.48145466, 0.4578275, 0.40821073], std=[0.26862954, 0.26130258, 0.27577711])

Dataset	EuroSAT
# classes	5
# train_x	80
# val	20
# test	3,900

Loading CLIP (backbone: ViT-B/16)
/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py:617: UserWarning: This DataLoader will create
warnings.warn(
/usr/local/lib/python3.10/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is deprecated
warnings.warn(
/content/ProMetaR/dassl/utils/torchtools.py:102: FutureWarning: You are using `torch.load` with `weights_only=False`
checkpoint = torch.load(fpath, map_location=map_location)
Building custom CLIP
Initial context: "a photo of a"
Number of context words (tokens): 4
Turning off gradients in both the image and the text encoder
Parameters to be updated: {'prompt_learner.meta_net.layer1.0.weight', 'prompt_learner.meta_net.layer3.0.weight',
Loading evaluator: Classification
Loading weights to prompt_learner from "outputs/cocoop/prompt_learner/model.pth.tar-100" (epoch = 100)
Evaluate on the *test* set
100%|██████████| 39/39 [00:59<00:00, 1.51s/it]=> result
* total: 3,900
* correct: 1,765
* accuracy: 45.3%
* error: 54.7%
* macro_f1: 38.9%