

✓ Preparation

```

!git clone https://github.com/mlvlab/ProMetaR.git
%cd ProMetaR/

!git clone https://github.com/KaiyangZhou/Dassl.pytorch.git
%cd Dassl.pytorch/

# Install dependencies
!pip install -r requirements.txt
!cp -r dassl ../
# Install this library (no need to re-build if the source code is modified)
# !python setup.py develop
%cd ..

!pip install -r requirements.txt

%mkdir outputs
%mkdir data

%cd data
%mkdir eurosat
!wget http://madm.dfki.de/files/sentinel/EuroSAT.zip EuroSAT.zip

!unzip -o EuroSAT.zip -d eurosat/
%cd eurosat
!gdown 1Ip7yaCWF0ea0FUGga0LUdVi_DDQth1o

%cd ../../

import os.path as osp
from collections import OrderedDict
import math
import torch
import torch.nn as nn
from torch.nn import functional as F
from torch.cuda.amp import GradScaler, autocast
from PIL import Image
import torchvision.transforms as transforms
import torch
from clip import clip
from clip.simple_tokenizer import SimpleTokenizer as _Tokenizer
import time
from tqdm import tqdm
import datetime
import argparse
from dassl.utils import setup_logger, set_random_seed, collect_env_info
from dassl.config import get_cfg_default
from dassl.engine import build_trainer
from dassl.engine import TRAINER_REGISTRY, TrainerX
from dassl.metrics import compute_accuracy
from dassl.utils import load_pretrained_weights, load_checkpoint
from dassl.optim import build_optimizer, build_lr_scheduler

# custom
import datasets.oxford_pets
import datasets.oxford_flowers
import datasets.fgvc_aircraft
import datasets.dtd
import datasets.eurosat
import datasets.stanford_cars
import datasets.food101
import datasets.sun397
import datasets.caltech101
import datasets.ucf101
import datasets.imagenet
import datasets.imagenet_sketch
import datasets.imagenetv2
import datasets.imagenet_a
import datasets.imagenet_r

def print_args(args, cfg):
    print("*****")
    print("** Arguments **")
    print("*****")
    optkeys = list(args.__dict__.keys())
    optkeys.sort()
    for key in optkeys:
        print("{}: {}".format(key, args.__dict__[key]))
    print("*****")
    print("** Config **")
    print("*****")
    print(cfg)

def reset_cfg(cfg, args):
    if args.root:
        cfg.DATASET.ROOT = args.root
    if args.output_dir:
        cfg.OUTPUT_DIR = args.output_dir
    if args.seed:
        cfg.SEED = args.seed
    if args.trainer:
        cfg.TRAINER.NAME = args.trainer
    cfg.DATASET.NUM_SHOTS = 16
    cfg.DATASET.SUBSAMPLE_CLASSES = args.subsample_classes

```

```

cfg.TRAINER.SUBSAMPLE_CLASSES = args.subsample_classes
cfg.DATALOADER.TRAIN_X.BATCH_SIZE = args.train_batch_size
cfg.OPTIM.MAX_EPOCH = args.epoch

def extend_cfg(cfg):
    """
    Add new config variables.
    """
    from yacs.config import CfgNode as CN
    cfg.TRAINER.COOP = CN()
    cfg.TRAINER.COOP.N_CTX = 16 # number of context vectors
    cfg.TRAINER.COOP.CSC = False # class-specific context
    cfg.TRAINER.COOP.CTX_INIT = "" # initialization words
    cfg.TRAINER.COOP.PREC = "fp16" # fp16, fp32, amp
    cfg.TRAINER.COOP.CLASS_TOKEN_POSITION = "end" # 'middle' or 'end' or 'front'
    cfg.TRAINER.COOCOOP = CN()
    cfg.TRAINER.COOCOOP.N_CTX = 4 # number of context vectors
    cfg.TRAINER.COOCOOP.CTX_INIT = "a photo of a" # initialization words
    cfg.TRAINER.COOCOOP.PREC = "fp16" # fp16, fp32, amp
    cfg.TRAINER.PROMETAR = CN()
    cfg.TRAINER.PROMETAR.N_CTX_VISION = 4 # number of context vectors at the vision branch
    cfg.TRAINER.PROMETAR.N_CTX_TEXT = 4 # number of context vectors at the language branch
    cfg.TRAINER.PROMETAR.CTX_INIT = "a photo of a" # initialization words
    cfg.TRAINER.PROMETAR.PREC = "fp16" # fp16, fp32, amp
    cfg.TRAINER.PROMETAR.PROMPT_DEPTH_VISION = 9 # Max 12, minimum 0, for 0 it will be using shallow IVLP prompting (J=1)
    cfg.TRAINER.PROMETAR.PROMPT_DEPTH_TEXT = 9 # Max 12, minimum 0, for 0 it will be using shallow IVLP prompting (J=1)
    cfg.DATASET.SUBSAMPLE_CLASSES = "all" # all, base or new
    cfg.TRAINER.PROMETAR.ADAPT_LR = 0.0005
    cfg.TRAINER.PROMETAR.LR_RATIO = 0.0005
    cfg.TRAINER.PROMETAR.FAST_ADAPTATION = False
    cfg.TRAINER.PROMETAR.MIXUP_ALPHA = 0.5
    cfg.TRAINER.PROMETAR.MIXUP_BETA = 0.5
    cfg.TRAINER.PROMETAR.DIM_RATE=8
    cfg.OPTIM_VNET = CN()
    cfg.OPTIM_VNET.NAME = "adam"
    cfg.OPTIM_VNET.LR = 0.0003
    cfg.OPTIM_VNET.WEIGHT_DECAY = 5e-4
    cfg.OPTIM_VNET.MOMENTUM = 0.9
    cfg.OPTIM_VNET.SGD_DAMPNING = 0
    cfg.OPTIM_VNET.SGD_NESTEROV = False
    cfg.OPTIM_VNET.RMSPROP_ALPHA = 0.99
    cfg.OPTIM_VNET.ADAM_BETA1 = 0.9
    cfg.OPTIM_VNET.ADAM_BETA2 = 0.999
    cfg.OPTIM_VNET.STAGED_LR = False
    cfg.OPTIM_VNET.NEW_LAYERS = ()
    cfg.OPTIM_VNET.BASE_LR_MULT = 0.1
    # Learning rate scheduler
    cfg.OPTIM_VNET.LR_SCHEDULER = "single_step"
    # -1 or 0 means the stepsize is equal to max_epoch
    cfg.OPTIM_VNET.STEPSIZE = (-1, )
    cfg.OPTIM_VNET.GAMMA = 0.1
    cfg.OPTIM_VNET.MAX_EPOCH = 10
    # Set WARMUP_EPOCH larger than 0 to activate warmup training
    cfg.OPTIM_VNET.WARMUP_EPOCH = -1
    # Either linear or constant
    cfg.OPTIM_VNET.WARMUP_TYPE = "linear"
    # Constant learning rate when type=constant
    cfg.OPTIM_VNET.WARMUP_CONS_LR = 1e-5
    # Minimum learning rate when type=linear
    cfg.OPTIM_VNET.WARMUP_MIN_LR = 1e-5
    # Recount epoch for the next scheduler (last_epoch=-1)
    # Otherwise last_epoch=warmup_epoch
    cfg.OPTIM_VNET.WARMUP_RECOUNT = True

def setup_cfg(args):
    cfg = get_cfg_default()
    extend_cfg(cfg)
    # 1. From the dataset config file
    if args.dataset_config_file:
        cfg.merge_from_file(args.dataset_config_file)
    # 2. From the method config file
    if args.config_file:
        cfg.merge_from_file(args.config_file)
    # 3. From input arguments
    reset_cfg(cfg, args)
    cfg.freeze()
    return cfg

_tokenizer = _Tokenizer()

def load_clip_to_cpu(cfg): # Load CLIP
    backbone_name = cfg.MODEL.BACKBONE.NAME
    url = clip._MODELS[backbone_name]
    model_path = clip._download(url)

    try:
        # loading JIT archive
        model = torch.jit.load(model_path, map_location="cpu").eval()
        state_dict = None

    except RuntimeError:
        state_dict = torch.load(model_path, map_location="cpu")

    if cfg.TRAINER.NAME == "":
        design_trainer = "CoOp"
    else:
        design_trainer = cfg.TRAINER.NAME
    design_details = {"trainer": design_trainer,
                     "vision_depth": 0,
                     "language_depth": 0,
                     "vision_ctx": 0,
                     "language_ctx": 0}

```

```

        language_ctx = v, vision_ctx = v,
        "language_ctx": 0}
    model = clip.build_model(state_dict or model.state_dict(), design_details)

    return model

from dassl.config import get_cfg_default
cfg = get_cfg_default()
cfg.MODEL.BACKBONE.NAME = "ViT-B/16" # Set the vision encoder backbone of CLIP to ViT.
clip_model = load_clip_to_cpu(cfg)

class TextEncoder(nn.Module):
    def __init__(self, clip_model): # 초기화 하는 함수
        super().__init__()
        self.transformer = clip_model.transformer
        self.positional_embedding = clip_model.positional_embedding
        self.ln_final = clip_model.ln_final
        self.text_projection = clip_model.text_projection
        self.dtype = clip_model.dtype

    def forward(self, prompts, tokenized_prompts): # 모델 호출
        x = prompts + self.positional_embedding.type(self.dtype)
        x = x.permute(1, 0, 2) # NLD -> LND
        x = self.transformer(x)
        x = x.permute(1, 0, 2) # LND -> NLD
        x = self.ln_final(x).type(self.dtype)

        # x.shape = [batch_size, n_ctx, transformer.width]
        # take features from the eot embedding (eot_token is the highest number in each sequence)
        x = x[torch.arange(x.shape[0]), tokenized_prompts.argmax(dim=-1)] @ self.text_projection

    return x

@TRAINER_REGISTRY.register(force=True)
class CoCoOp(TrainerX):
    def check_cfg(self, cfg):
        assert cfg.TRAINER.COCOOP.PREC in ["fp16", "fp32", "amp"]

    def build_model(self):
        cfg = self.cfg
        classnames = self.dm.dataset.classnames
        print(f>Loading CLIP (backbone: {cfg.MODEL.BACKBONE.NAME})")
        clip_model = load_clip_to_cpu(cfg)

        if cfg.TRAINER.COCOOP.PREC == "fp32" or cfg.TRAINER.COCOOP.PREC == "amp":
            # CLIP's default precision is fp16
            clip_model.float()

        print("Building custom CLIP")
        self.model = CoCoOpCustomCLIP(cfg, classnames, clip_model)

        print("Turning off gradients in both the image and the text encoder")
        name_to_update = "prompt_learner"

        for name, param in self.model.named_parameters():
            if name_to_update not in name:
                param.requires_grad_(False)

        # Double check
        enabled = set()
        for name, param in self.model.named_parameters():
            if param.requires_grad:
                enabled.add(name)
        print(f>Parameters to be updated: {enabled}")

        if cfg.MODEL.INIT_WEIGHTS:
            load_pretrained_weights(self.model.prompt_learner, cfg.MODEL.INIT_WEIGHTS)

        self.model.to(self.device)
        # NOTE: only give prompt_learner to the optimizer
        self.optim = build_optimizer(self.model.prompt_learner, cfg.OPTIM)
        self.sched = build_lr_scheduler(self.optim, cfg.OPTIM)
        self.register_model("prompt_learner", self.model.prompt_learner, self.optim, self.sched)

        self.scaler = GradScaler() if cfg.TRAINER.COCOOP.PREC == "amp" else None

        # Note that multi-gpu training could be slow because CLIP's size is
        # big, which slows down the copy operation in DataParallel
        device_count = torch.cuda.device_count()
        if device_count > 1:
            print(f>Multiple GPUs detected (n_gpus={device_count}), use all of them!")
            self.model = nn.DataParallel(self.model)

    def before_train(self):
        directory = self.cfg.OUTPUT_DIR
        if self.cfg.RESUME:
            directory = self.cfg.RESUME
        self.start_epoch = self.resume_model_if_exist(directory)

        # Remember the starting time (for computing the elapsed time)
        self.time_start = time.time()

    def forward_backward(self, batch):
        image, label = self.parse_batch_train(batch)

        model = self.model

```

```

model = self.model
optim = self.optim
scaler = self.scaler

prec = self.cfg.TRAINER.COOCOOP.PREC
loss = model(image, label) # Input image 모델 통과
optim.zero_grad()
loss.backward() # Backward (역전파)
optim.step() # 모델 parameter update

loss_summary = {"loss": loss.item()}

if (self.batch_idx + 1) == self.num_batches:
    self.update_lr()

return loss_summary

def parse_batch_train(self, batch):
    input = batch["img"]
    label = batch["label"]
    input = input.to(self.device)
    label = label.to(self.device)
    return input, label

def load_model(self, directory, epoch=None):
    if not directory:
        print("Note that load_model() is skipped as no pretrained model is given")
        return

    names = self.get_model_names()

    # By default, the best model is loaded
    model_file = "model-best.pth.tar"

    if epoch is not None:
        model_file = "model.pth.tar-" + str(epoch)

    for name in names:
        model_path = osp.join(directory, name, model_file)

        if not osp.exists(model_path):
            raise FileNotFoundError('Model not found at "{}".format(model_path))

        checkpoint = load_checkpoint(model_path)
        state_dict = checkpoint["state_dict"]
        epoch = checkpoint["epoch"]

        # Ignore fixed token vectors
        if "token_prefix" in state_dict:
            del state_dict["token_prefix"]

        if "token_suffix" in state_dict:
            del state_dict["token_suffix"]

        print("Loading weights to {} " 'from "{}' (epoch = {})'.format(name, model_path, epoch))
        # set strict=False
        self._models[name].load_state_dict(state_dict, strict=False)

def after_train(self):
    print("Finish training")

    do_test = not self.cfg.TEST.NO_TEST
    if do_test:
        if self.cfg.TEST.FINAL_MODEL == "best_val":
            print("Deploy the model with the best val performance")
            self.load_model(self.output_dir)
        else:
            print("Deploy the last-epoch model")
            acc = self.test()

    # Show elapsed time
    elapsed = round(time.time() - self.time_start)
    elapsed = str(datetime.timedelta(seconds=elapsed))
    print(f"Elapsed: {elapsed}")

    # Close writer
    self.close_writer()
    return acc

def train(self):
    """Generic training loops."""
    self.before_train()
    for self.epoch in range(self.start_epoch, self.max_epoch):
        self.before_epoch()
        self.run_epoch()
        self.after_epoch()
    acc = self.after_train()
    return acc

parser = argparse.ArgumentParser()
parser.add_argument("--root", type=str, default="data/", help="path to dataset")
parser.add_argument("--output-dir", type=str, default="outputs/cocoop3", help="output directory")
parser.add_argument(
    "--seed", type=int, default=1, help="only positive value enables a fixed seed"
)
parser.add_argument(
    "--config-file", type=str, default="configs/trainers/ProMetaR/vit_b16_c2_ep10_batch4_4+4ctx.yaml", help="path to config file"
)
parser.add_argument(
    "--data-path", type=str, default="data/", help="path to dataset"
)

```

```
    "--dataset-config-file",
    type=str,
    default="configs/datasets/eurosat.yaml",
    help="path to config file for dataset setup",
)
parser.add_argument("--trainer", type=str, default="CoOp", help="name of trainer")
parser.add_argument("--eval-only", action="store_true", help="evaluation only")
parser.add_argument(
    "--model-dir",
    type=str,
    default="",
    help="load model from this directory for eval-only mode",
)
parser.add_argument("--train-batch-size", type=int, default=4)
parser.add_argument("--epoch", type=int, default=10)
parser.add_argument("--subsample-classes", type=str, default="base")
parser.add_argument(
    "--load-epoch", type=int, default=0, help="load model weights at this epoch for evaluation"
)
args = parser.parse_args([])

def main(args):
    cfg = setup_cfg(args)
    if cfg.SEED >= 0:
        set_random_seed(cfg.SEED)

    if torch.cuda.is_available() and cfg.USE_CUDA:
        torch.backends.cudnn.benchmark = True

    trainer = build_trainer(cfg)
    if args.eval_only:
        trainer.load_model(args.model_dir, epoch=args.load_epoch)
        acc = trainer.test()
        return acc

    acc = trainer.train()
    return acc
```



Implementing CoCoOp

- <https://colab.research.google.com/drive/1DkZ8TiVACVtRFWxhp1WpzVEJBEBPkzQW#scrollTo=1lKdgiKFsoj&printMode=true>

```

        [
            prefix, # (dim0, 1, dim)
            ctx, # (dim0, n_ctx, dim)
            suffix, # (dim0, *, dim)
        ],
        dim=1,
    )

    return prompts

def forward(self, im_features):
    prefix = self.token_prefix
    suffix = self.token_suffix
    ctx = self.ctx # (n_ctx, ctx_dim)

    #####
    #Image feature is given as input to meta network # (batch, ctx_dim)
    bias = self.meta_net(im_features) # (batch, ctx_dim)
    bias = bias.unsqueeze(1) # (batch, 1, ctx_dim)
    ctx = ctx.unsqueeze(0) # (1, n_ctx, ctx_dim)
    #ctx_shifted = ctx + " Add meta token to context token" # (batch, n_ctx, ctx_dim)
    ctx_shifted = ctx + bias # (batch, n_ctx, ctx_dim)
    #####
    #####

    # Use instance-conditioned context tokens for all classes
    prompts = []
    for ctx_shifted_i in ctx_shifted:
        ctx_i = ctx_shifted_i.unsqueeze(0).expand(self.n_cls, -1, -1)
        pts_i = self.construct_prompts(ctx_i, prefix, suffix) # (n_cls, n_tkn, ctx_dim)
        prompts.append(pts_i)
    prompts = torch.stack(prompts)

    return prompts

class CoCoOpCustomCLIP(nn.Module):
    def __init__(self, cfg, classnames, clip_model):
        super().__init__()
        self.prompt_learner = CoCoOpPromptLearner(cfg, classnames, clip_model)
        self.tokenized_prompts = self.prompt_learner.tokenized_prompts
        self.image_encoder = clip_model.visual
        self.text_encoder = TextEncoder(clip_model)
        self.logit_scale = clip_model.logit_scale
        self.dtype = clip_model.dtype

    def forward(self, image, label=None):
        tokenized_prompts = self.tokenized_prompts
        logit_scale = self.logit_scale.exp()

        image_features = self.image_encoder(image.type(self.dtype))
        image_features = image_features / image_features.norm(dim=-1, keepdim=True)

        #####
        prompts = self.prompt_learner(image_features)
        #####

        logits = []
        for pts_i, imf_i in zip(prompts, image_features):
            text_features = self.text_encoder(pts_i, tokenized_prompts)
            text_features = text_features / text_features.norm(dim=-1, keepdim=True)
            l_i = logit_scale * imf_i @ text_features.t()
            logits.append(l_i)
        logits = torch.stack(logits)

        if self.prompt_learner.training:
            return F.cross_entropy(logits, label)

        return logits

```

▼ Training CoCoOp

- Train CoCoOp on the EuroSAT dataset.

```
!pwd
!ls
```



```

/content/ProMetaR
clip          dassl          datasets      meta_learning  README.md      trainers
clip_words.csv Dassl.pytorch docs          outputs        requirements.txt train.py
configs       data          LICENSE      parse_test_res.py scripts

```

```

# Train on the Base Classes Train split and evaluate accuracy on the Base Classes Test split.
args.trainer = "CoCoOp"
args.train_batch_size = 4
args.epoch = 100

```

```
args.output_dir = "outputs/cocoop"
```

```
args.subsample_classes = "base"
```

```
args.eval_only = False
```

```
cocoop_base_acc = main(args)
```

```

→ epoch [55/100] batch [20/20] time 0.093 (0.127) data 0.000 (0.018) loss 0.2917 (0.1799) lr 1.0933e-03 eta 0:01:54
epoch [56/100] batch [20/20] time 0.092 (0.127) data 0.000 (0.018) loss 0.2384 (0.2613) lr 1.0545e-03 eta 0:01:51
epoch [57/100] batch [20/20] time 0.135 (0.155) data 0.000 (0.023) loss 0.3364 (0.3352) lr 1.0158e-03 eta 0:02:12
epoch [58/100] batch [20/20] time 0.101 (0.187) data 0.000 (0.034) loss 0.3237 (0.2660) lr 9.7732e-04 eta 0:02:37
epoch [59/100] batch [20/20] time 0.092 (0.128) data 0.000 (0.019) loss 0.0295 (0.2851) lr 9.3914e-04 eta 0:01:44
epoch [60/100] batch [20/20] time 0.103 (0.128) data 0.000 (0.019) loss 0.0961 (0.1896) lr 9.0126e-04 eta 0:01:42
epoch [61/100] batch [20/20] time 0.094 (0.128) data 0.000 (0.022) loss 0.3149 (0.2265) lr 8.6373e-04 eta 0:01:39
epoch [62/100] batch [20/20] time 0.144 (0.168) data 0.000 (0.018) loss 0.0041 (0.2124) lr 8.2658e-04 eta 0:02:07
epoch [63/100] batch [20/20] time 0.094 (0.127) data 0.000 (0.018) loss 0.1748 (0.2624) lr 7.8984e-04 eta 0:01:34
epoch [64/100] batch [20/20] time 0.094 (0.131) data 0.000 (0.020) loss 0.2600 (0.1714) lr 7.5357e-04 eta 0:01:34
epoch [65/100] batch [20/20] time 0.097 (0.129) data 0.000 (0.021) loss 0.5747 (0.2100) lr 7.1778e-04 eta 0:01:30
epoch [66/100] batch [20/20] time 0.151 (0.155) data 0.000 (0.025) loss 0.1279 (0.1686) lr 6.8251e-04 eta 0:01:45
epoch [67/100] batch [20/20] time 0.136 (0.195) data 0.000 (0.034) loss 0.0054 (0.2219) lr 6.4781e-04 eta 0:02:08
epoch [68/100] batch [20/20] time 0.095 (0.127) data 0.000 (0.020) loss 0.2773 (0.2684) lr 6.1370e-04 eta 0:01:21
epoch [69/100] batch [20/20] time 0.096 (0.130) data 0.000 (0.023) loss 0.0228 (0.2471) lr 5.8022e-04 eta 0:01:20
epoch [70/100] batch [20/20] time 0.095 (0.128) data 0.000 (0.021) loss 0.2318 (0.1503) lr 5.4740e-04 eta 0:01:16
epoch [71/100] batch [20/20] time 0.161 (0.149) data 0.000 (0.023) loss 0.0285 (0.1188) lr 5.1527e-04 eta 0:01:26
epoch [72/100] batch [20/20] time 0.140 (0.196) data 0.000 (0.034) loss 0.1163 (0.2144) lr 4.8387e-04 eta 0:01:49
epoch [73/100] batch [20/20] time 0.095 (0.137) data 0.000 (0.019) loss 0.0424 (0.1745) lr 4.5322e-04 eta 0:01:13
epoch [74/100] batch [20/20] time 0.096 (0.129) data 0.000 (0.026) loss 0.1774 (0.1305) lr 4.2336e-04 eta 0:01:07
epoch [75/100] batch [20/20] time 0.098 (0.128) data 0.000 (0.020) loss 0.0523 (0.1880) lr 3.9432e-04 eta 0:01:03
epoch [76/100] batch [20/20] time 0.144 (0.155) data 0.000 (0.018) loss 0.0109 (0.1781) lr 3.6612e-04 eta 0:01:14
epoch [77/100] batch [20/20] time 0.092 (0.182) data 0.000 (0.033) loss 0.0092 (0.1832) lr 3.3879e-04 eta 0:01:23
epoch [78/100] batch [20/20] time 0.099 (0.130) data 0.000 (0.023) loss 0.1420 (0.2149) lr 3.1236e-04 eta 0:00:57
epoch [79/100] batch [20/20] time 0.093 (0.127) data 0.000 (0.019) loss 0.6455 (0.2502) lr 2.8686e-04 eta 0:00:53
epoch [80/100] batch [20/20] time 0.106 (0.127) data 0.000 (0.021) loss 0.1262 (0.1671) lr 2.6231e-04 eta 0:00:50
epoch [81/100] batch [20/20] time 0.154 (0.170) data 0.000 (0.020) loss 0.1049 (0.1736) lr 2.3873e-04 eta 0:01:04
epoch [82/100] batch [20/20] time 0.095 (0.130) data 0.000 (0.020) loss 0.5278 (0.1947) lr 2.1615e-04 eta 0:00:46
epoch [83/100] batch [20/20] time 0.093 (0.136) data 0.000 (0.026) loss 0.1053 (0.1895) lr 1.9459e-04 eta 0:00:46
epoch [84/100] batch [20/20] time 0.105 (0.130) data 0.000 (0.020) loss 0.1261 (0.1526) lr 1.7407e-04 eta 0:00:41
epoch [85/100] batch [20/20] time 0.138 (0.154) data 0.000 (0.019) loss 0.0314 (0.1640) lr 1.5462e-04 eta 0:00:46
epoch [86/100] batch [20/20] time 0.159 (0.193) data 0.000 (0.029) loss 0.0459 (0.1491) lr 1.3624e-04 eta 0:00:54
epoch [87/100] batch [20/20] time 0.116 (0.129) data 0.000 (0.022) loss 0.2108 (0.1862) lr 1.1897e-04 eta 0:00:33
epoch [88/100] batch [20/20] time 0.093 (0.129) data 0.000 (0.019) loss 0.1178 (0.2581) lr 1.0281e-04 eta 0:00:30
epoch [89/100] batch [20/20] time 0.093 (0.128) data 0.000 (0.019) loss 0.0460 (0.2158) lr 8.7779e-05 eta 0:00:28
epoch [90/100] batch [20/20] time 0.142 (0.163) data 0.000 (0.021) loss 0.0492 (0.1039) lr 7.3899e-05 eta 0:00:32
epoch [91/100] batch [20/20] time 0.100 (0.134) data 0.000 (0.027) loss 0.2791 (0.1459) lr 6.1179e-05 eta 0:00:24
epoch [92/100] batch [20/20] time 0.098 (0.128) data 0.000 (0.024) loss 0.0514 (0.1019) lr 4.9633e-05 eta 0:00:20
epoch [93/100] batch [20/20] time 0.093 (0.138) data 0.000 (0.023) loss 0.1763 (0.2449) lr 3.9271e-05 eta 0:00:19
epoch [94/100] batch [20/20] time 0.122 (0.145) data 0.000 (0.020) loss 0.2859 (0.2261) lr 3.0104e-05 eta 0:00:17
epoch [95/100] batch [20/20] time 0.139 (0.197) data 0.000 (0.035) loss 0.1564 (0.1853) lr 2.2141e-05 eta 0:00:19
epoch [96/100] batch [20/20] time 0.095 (0.129) data 0.000 (0.019) loss 0.4089 (0.1330) lr 1.5390e-05 eta 0:00:10
epoch [97/100] batch [20/20] time 0.095 (0.129) data 0.000 (0.021) loss 0.0698 (0.1542) lr 9.8566e-06 eta 0:00:07
epoch [98/100] batch [20/20] time 0.095 (0.129) data 0.000 (0.023) loss 0.2188 (0.2041) lr 5.5475e-06 eta 0:00:05
epoch [99/100] batch [20/20] time 0.157 (0.150) data 0.000 (0.019) loss 0.0691 (0.1264) lr 2.4666e-06 eta 0:00:03
epoch [100/100] batch [20/20] time 0.151 (0.201) data 0.000 (0.038) loss 0.0025 (0.1101) lr 6.1680e-07 eta 0:00:00
Checkpoint saved to outputs/cocoop/prompt_learner/model.pth.tar-100
Finish training
Deploy the last-epoch model
Evaluate on the *test* set
100%|██████████| 42/42 [01:06<00:00, 1.57s/it]=> result
* total: 4,200
* correct: 3,813
* accuracy: 90.8%
* error: 9.2%
* macro_f1: 90.9%
Elapsed: 0:06:33

```

```
# Accuracy on the New Classes.
```

```
args.model_dir = "outputs/cocoop"
```

```
args.output_dir = "outputs/cocoop/new_classes"
```

```
args.subsample_classes = "new"
```

```
args.load_epoch = 100
```

```
args.eval_only = True
```

```
coop_novel_acc = main(args)
```

```

→ Loading trainer: CoCoOp
Loading dataset: EuroSAT
Reading split from /content/ProMetaR/data/eurosat/split_zhou_EuroSAT.json
Loading preprocessed few-shot data from /content/ProMetaR/data/eurosat/split_fewshot/shot_16-seed_1.pkl
SUBSAMPLE NEW CLASSES!
Building transform_train
+ random resized crop (size=(224, 224), scale=(0.08, 1.0))
+ random flip
+ to torch tensor of range [0, 1]
+ normalization (mean=[0.48145466, 0.4578275, 0.40821073], std=[0.26862954, 0.26130258, 0.27577711])
Building transform_test
+ resize the smaller edge to 224
+ 224x224 center crop

```



```

+ to torch tensor of range [0, 1]
+ normalization (mean=[0.48145466, 0.4578275, 0.40821073], std=[0.26862954, 0.26130258, 0.27577711])
-----
Dataset      EuroSAT
# classes    5
# train_x    80
# val        20
# test       3,900
-----
Loading CLIP (backbone: ViT-B/16)
/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py:617: UserWarning: This DataLoader will cre
warnings.warn(
/usr/local/lib/python3.10/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is dep
warnings.warn(
/content/ProMetaR/dassl/utils/torchtools.py:102: FutureWarning: You are using `torch.load` with `weights_only=False`
checkpoint = torch.load(fpath, map_location=map_location)
Building custom CLIP
Initial context: "a photo of a"
Number of context words (tokens): 4
Turning off gradients in both the image and the text encoder
Parameters to be updated: {'prompt_learner.meta_net.linear2.weight', 'prompt_learner.meta_net.linear1.bias', 'pro
Loading evaluator: Classification
Loading weights to prompt_learner from "outputs/cocoop/prompt_learner/model.pth.tar-100" (epoch = 100)
Evaluate on the *test* set
100%|██████████| 39/39 [01:01<00:00, 1.58s/it]=> result
* total: 3,900
* correct: 1,687
* accuracy: 43.3%
* error: 56.7%
* macro_f1: 39.0%

```

Next Step: Tries to improve CoCoOp

- current meta network is a simple two-layer MLP with ReLU activation.
- try other architectures (CoCoOp Meta Network Variants)

1. Deeper Network (variant='deeper')

- Adds more layers and depth
- Includes dropout for regularization
- Uses a gradual dimension reduction

2. Residual Network (variant='residual')

- Adds skip connections
- Better gradient flow
- Might help with training stability

3. Transformer-based (variant='transformer')

- Uses self-attention mechanism
- Could better capture complex relationships in visual features
- Includes adaptive pooling

(See another notebooks)

